

# ML-Enhanced Large Neighborhood Search Project Report

Machine Learning for Optimization, WS 2025

Daniel Levin (12433760)

## 1 Problem: SCF-PDP

The **Selective Capacitated Fair Pickup and Delivery Problem (SCF-PDP)** is a vehicle routing problem where the goal is to design fair and feasible routes for a subset of customer requests. Each customer needs transportation of goods from a pickup location to a corresponding drop-off location.

### 1.1 Problem Formulation

The problem is modeled on a complete directed graph  $G = (V, A)$  where:

- $V$  contains the vehicle depot plus all pickup/drop-off locations
- $A = \{(u, v) : u, v \in V, u \neq v\}$  represents travel routes
- Arc distances  $a_{u,v} = \lceil \sqrt{(x_u - x_v)^2 + (y_u - y_v)^2} \rceil$  (rounded Euclidean)

There are  $n$  customer requests  $CR = \{1, \dots, n\}$ , each defined by pickup location  $v_i^\uparrow$  and drop-off location  $v_i^\downarrow$  with demand  $c_i$ . A fleet  $K$  of  $n_K$  identical vehicles with capacity  $C$  serves requests starting from the depot.

### 1.2 Constraints and Objective

A feasible solution assigns routes  $R_k$  to each vehicle  $k \in K$  such that:

- Vehicle capacity never exceeded along any route
- Each served request handled entirely by one vehicle
- At least  $\gamma$  requests served across all vehicles

The objective minimizes total distance plus fairness penalty:

$$\sum_{k \in K} d(R_k) + \rho \cdot (1 - J(R))$$

where  $d(R_k)$  is route  $k$ 's total distance and  $J(R)$  is the Jain fairness index:

$$J(R) = \frac{\left(\sum_{k \in K} d(R_k)\right)^2}{n_K \cdot \sum_{k \in K} d(R_k)^2}$$

The fairness index  $J(R) \in (0, 1]$  equals 1 when all routes have equal distance. Parameter  $\rho$  controls the distance-fairness trade-off.

## 2 Baseline: Tuned ALNS

The baseline implementation uses Adaptive Large Neighborhood Search (ALNS) with simulated annealing acceptance, developed and tuned in the Heuristic Optimization Techniques course. This section summarizes the key components.

### 2.1 Operators

The implementation uses 3 destroy and 3 repair operators, yielding 9 combinations. Each destroy operator removes  $q$  requests sampled uniformly from [10%, 40%] of served requests.

#### Destroy Operators:

- **Random**: Removes  $q$  randomly selected requests (unbiased exploration)
- **WorstCost**: Removes requests with highest distance contribution
- **LongestRoute**: Removes from longest routes (targets fairness)

#### Repair Operators:

- **Greedy**: Reinserts using flexible pickup-dropoff construction heuristic
- **RandomGreedy**: Reinserts at random feasible positions
- **ObjectiveAware**: Minimizes full objective including fairness penalty

### 2.2 Adaptive Weight Mechanism

Operators are selected via roulette wheel with adaptive weights. Scores: 10 for new best, 1 for accepted, 0 for rejected. Weights update every 100 iterations using:

$$\rho_i \leftarrow \rho_i \cdot (1 - \gamma) + \gamma \cdot \frac{s_i}{a_i}$$

where  $\gamma$  is the reaction factor,  $s_i$  total score, and  $a_i$  applications in the period.

### 2.3 Parameters

Parameter	Description	Default
<code>max_iterations</code>	Maximum iterations	10000
<code>max_time_seconds</code>	Runtime limit	300s
<code>weight_update_period</code>	Weight update frequency	100
<code>reaction_factor (<math>\gamma</math>)</code>	Adaptation speed	0.1
<code>min/max_removal_pct</code>	Removal range	10%-40%
<code>initial_temperature</code>	SA start temperature	100.0
<code>cooling_rate</code>	Temperature decay	0.99

Table 1: ALNS default parameters

### **3 Multi-Armed Bandit Operator Selection**

3.1 Method

3.2 Results

### **4 Reinforcement Learning Operator Selection**

4.1 Method

4.2 Results

### **5 Comparison and Analysis**

### **6 Conclusion**